

Dec 06, 15 2:06

sched.h

Page 1/1

```

#ifndef _SCHED_H_
#define _SCHED_H_

#include "context.h"

#define STACK_SIZE 0xffff

#define SCHED_NPROC      512
#define SCHED_READY     0
#define SCHED_RUNNING   1
#define SCHED_SLEEPING  2
#define SCHED_ZOMBIE    3

struct sched_proc {
    int state;
    int pid;
    int ppid;
    int nice;
    unsigned accumulated;
    unsigned cpu_time;
    void *stack;
    struct savctx context;
    int exit_code;
    int priority;
};

struct sched_waitq {
    struct sched_proc * procs [SCHED_NPROC];
    unsigned nprocs;
};

int create_pid();
void * create_stack();
struct sched_proc * create_proc(int ppid);
int sched_init(void (*init_fn)());
int sched_fork();
void sched_exit(int code);
int sched_wait(int *exit_code);
void sched_nice(int val);
int sched_getpid();
int sched_getppid();
int sched_gettick();
void sched_ps();
int getPriority(int pid);
void sched_switch();
void sched_tick();

#endif /* _SCHED_H_ */

```

Dec 06, 15 23:08

sched.c

Page 1/5

```

#include "sched.h"
#include "adjstack.h"
#include "context.h"

#include <sys/time.h>
#include <sys/mman.h>
#include <stdlib.h>
#include <stdio.h>
#include <errno.h>
#include <signal.h>
#include <string.h>

#define HOLD_SIGNALS(x) \
    sigprocmask(SIG_BLOCK, &all_signals, NULL); \
    x \
    sigprocmask(SIG_UNBLOCK, &all_signals, NULL);

#define BLOCK_SIGNALS() sigprocmask(SIG_BLOCK, &all_signals, NULL)
#define UNBLOCK_SIGNALS() sigprocmask(SIG_UNBLOCK, &all_signals, NULL)

struct sched_proc * current = NULL;
struct sched_waitq * queue = NULL;
unsigned nticks;
sigset_t all_signals;

int create_pid() {
    if (queue == NULL) {
        fprintf(stderr, "Running process list has not been allocated!");
        return -1;
    }
    // Skip PID 0, since scheduler should probably have that
    for (unsigned i = 1; i < SCHED_NPROC; i++){
        if (queue->procs[i] == NULL)
            return i;
    }
    fprintf(stderr, "Ran out of PIDs!\n");
    return -1;
}

void * create_stack() {
    void * newsp;
    if ((newsp = mmap(0, STACK_SIZE, PROT_READ | PROT_WRITE, MAP_PRIVATE | MAP_A
NONYMOUS, 0, 0)) == MAP_FAILED) {
        perror("mmap()");
        return MAP_FAILED;
    }
    return newsp;
}

struct sched_proc * create_proc(int ppid) {
    int pid = create_pid();
    if (pid == -1) {
        return 0;
    }
    struct sched_proc * proc = (struct sched_proc *) malloc(sizeof(struct sched_
proc));
    if (proc == NULL) {
        perror("malloc()");
    }

```

Dec 06, 15 23:08

sched.c

Page 2/5

```

        return NULL;
    }
    void * newsp = create_stack();
    if (newsp == -1) {
        free(proc);
        return NULL;
    }

    // Set contents
    proc->state      = SCHED_READY;
    proc->pid        = pid;
    proc->ppid       = ppid;
    proc->nice        = 0;
    proc->accumulated = 0;
    proc->cpu_time    = 0;
    proc->stack       = newsp;

    queue->procs[pid] = proc;
    queue->nprocs++;

    return proc;
}

int sched_init(void (*init_fn)()) {
    // Initialize process queue
    queue = (struct sched_waitq *) calloc(1, sizeof(struct sched_waitq));
    if (queue == NULL) {
        perror("malloc()");
        return -1;
    }

    // Set up ticks
    sigfillset(&all_signals);
    nticks = 0;
    struct timeval t;
    t.tv_sec = 0;
    t.tv_usec = 1e5;
    struct itimerval it;
    it.it_interval = t;
    it.it_value = t;
    if (setitimer(ITIMER_VIRTUAL, &it, NULL) == -1) {
        perror("Failed to set timer");
        return -1;
    }
    signal(SIGVTALRM, sched_tick);

    // Create init process and switch control
    current = create_proc(0);
    struct savectx ctx;
    ctx.regs[JB_BP] = current->stack + STACK_SIZE;
    ctx.regs[JB_SP] = current->stack + STACK_SIZE;
    ctx.regs[JB_PC] = init_fn;
    restorectx(&ctx, current->pid);

    return 0;
}

int sched_fork() {
    int ret;

```

Dec 06, 15 23:08

sched.c

Page 3/5

```

    HOLD_SIGNALS(
        struct sched_proc * newproc = create_proc(current->pid);
        newproc->cpu_time = current->cpu_time;
        memcpy(newproc->stack, current->stack, STACK_SIZE);
        adjstack(newproc->stack, newproc->stack + STACK_SIZE, newproc->stack - c
current->stack);

        ret = savectx(&newproc->context);
        if (ret == 0) {
            newproc->context.regs[JB_BP] += newproc->stack - current->stack;
            newproc->context.regs[JB_SP] += newproc->stack - current->stack;
        }

    );
    if (ret == 0) {
        return newproc->pid;
    } else {
        current = queue->procs[ret];
        return 0;
    }
}

void sched_exit(int code) {
    HOLD_SIGNALS(
        current->state = SCHED_ZOMBIE;
        current->exit_code = code;
        queue->nprocs--;

        // Check for wait()
        if ((queue->procs[current->ppid] != NULL)
            && (queue->procs[current->ppid]->state == SCHED_SLEEPING)) {
            queue->procs[current->ppid]->state = SCHED_READY;
        }
    );
    sched_switch();
}

int sched_wait(int *exit_code) {
    while(1) {
        BLOCK_SIGNALS();
        int found = 0;
        int zombie = 0;
        found = 0;
        zombie = 0;
        for (int i = 1; i < SCHED_NPROC; i++) {
            if (queue->procs[i] != NULL
                && queue->procs[i]->ppid == current->pid) {
                found = 1;
                if (queue->procs[i]->state == SCHED_ZOMBIE) {
                    zombie = 1;
                    current->state = SCHED_READY;
                    *exit_code = queue->procs[i]->exit_code;
                    free(queue->procs[i]);
                    queue->procs[i] = NULL;
                    break;
                }
            }
        }
    }
}

```

Dec 06, 15 23:08

sched.c

Page 4/5

```

        if (found == 0) {
            UNBLOCK_SIGNALS();
            return -1;
        } else if (found == 1 && zombie == 1) {
            UNBLOCK_SIGNALS();
            return 0;
        } else {
            UNBLOCK_SIGNALS();
            current->state = SCHED_SLEEPING;
            sched_switch();
        }
    }
}

void sched_nice(int val) {
    val = val > 19 ? 19 : val;
    val = val < -20 ? -20 : val;
    current->nice = val;
}

int sched_getpid() {
    return current->pid;
}

int sched_getppid() {
    return current->ppid;
}

int sched_gettick() {
    return nticks;
}

void sched_ps() {
    printf("PID    PPID    STATE    STACK    STATIC    DYNAMIC    TIME    \n");
    for (int i = 1; i < SCHED_NPROC; i++) {
        if (queue->procs[i] != NULL) {
            printf("%-9d %-9d %-9d %-9x %-9d %-9d %-9d\n",
                queue->procs[i]->pid,
                queue->procs[i]->ppid,
                queue->procs[i]->state,
                (unsigned)queue->procs[i]->stack,
                20 - queue->procs[i]->nice,
                queue->procs[i]->priority,
                queue->procs[i]->accumulated);
        }
    }
}

int getPriority(int pid) {
    if (queue->procs[pid] != NULL) {
        int priority = 20 - queue->procs[pid]->nice - (int)queue->procs[pid]->accumulated/2;
        // priority = priority > 39 ? 39 : priority;
        // priority = priority < 0 ? 0 : priority;
        queue->procs[pid]->priority = priority;
        return priority;
    } else {

```

Dec 06, 15 23:08

sched.c

Page 5/5

```

        return -1;
    }
}

void sched_switch() {
    BLOCK_SIGNALS();
    if (current->state != SCHED_ZOMBIE && current->state != SCHED_SLEEPING)
        current->state = SCHED_READY;

    int best = -1;
    int best_pid = -1;
    int p;
    for (int pid = 1; pid < SCHED_NPROC; pid++) {
        if ((queue->procs[pid] != NULL)
            && (queue->procs[pid]->state == SCHED_READY)
            && ((p = getPriority(pid)) > best)) {
            best = p;
            best_pid = pid;
        }
    }

    if (best_pid == -1)
        return;

    if (best_pid == current->pid) {
        current->state = SCHED_RUNNING;
        current->cpu_time = 0;
        UNBLOCK_SIGNALS();
        return;
    }

    if (savectx(&current->context) == 0) {
        current = queue->procs[best_pid];
        current->cpu_time = 0;
        current->state = SCHED_RUNNING;
        UNBLOCK_SIGNALS();
        restorectx(&current->context, current->pid);
    }
}

void sched_tick() {
    sched_ps();
    nticks++;
    current->accumulated++;
    current->cpu_time++;
    sched_switch();
    return;
}

```

Dec 06, 15 23:13

sched_test.c

Page 1/2

```

#include "sched.h"

#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>

void child() {
    // priority: 6, 4, 5, 3, 2
    switch (sched_getpid()) {
        case 2:
            sched_nice(10);
            break;
        case 3:
            sched_nice(8);
            break;
        case 4:
            sched_nice(4);
            break;
        case 5:
            sched_nice(6);
            break;
        case 6:
            sched_nice(2);
            break;
    }
    printf("In child with pid %d, ppid %d\n", sched_getpid(), sched_getppid());
    unsigned long max = 1e8;
    for (unsigned long i = 0; i < max; i++) {
        getpid();
    }
    sched_exit(sched_getpid());
}

void init() {
    printf("Hello\n");
    int pid;
    for (int i = 0; i < 5; i++) {
        switch(pid = sched_fork()) {
            case -1:
                printf("sched_fork() error\n");
                sched_exit(0);
            case 0:
                child();
            default:
                printf("Created pid %d\n", pid);
                break;
        }
    }
    printf("In parent with pid %d\n", sched_getpid());
    int code;
    // unsigned long max = 1e9;
    // for (unsigned long i = 0; i < max; i++) {
    //     getpid();
    // }
    int order[5];
    for (int i = 0; i < 5; i++) {
        sched_wait(&code);
        order[i] = code;
    }
}

```

Dec 06, 15 23:13

sched_test.c

Page 2/2

```
    printf("Child returned with code %d\n", code);
}
printf("Children returned in order:");
for (int i = 0; i < 5; i++)
    printf(" %d", order[i]);
printf("\n");
sched_exit(0);
printf("Ending init process...\n");
exit(0);
}

int main(int argc, char **argv) {
    sched_init(init);
}
```


Dec 09, 15 0:39

out.txt

Page 1/3

Hello

Created pid 2

Created pid 3

Created pid 4

Created pid 5

Created pid 6

In parent with pid 1

In child with pid 2, ppid 1

PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	1	49f7a000	10	20	0
3	1	0	49f6a000	20	20	0
4	1	0	49f5a000	20	20	0
5	1	0	49f4a000	20	20	0
6	1	0	49f3a000	20	20	0

In child with pid 3, ppid 1

PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	10	1
3	1	1	49f6a000	12	20	0
4	1	0	49f5a000	20	20	0
5	1	0	49f4a000	20	20	0
6	1	0	49f3a000	20	20	0

In child with pid 4, ppid 1

PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	10	1
3	1	0	49f6a000	12	12	1
4	1	1	49f5a000	16	20	0
5	1	0	49f4a000	20	20	0
6	1	0	49f3a000	20	20	0

In child with pid 5, ppid 1

PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	10	1
3	1	0	49f6a000	12	12	1
4	1	0	49f5a000	16	16	1
5	1	1	49f4a000	14	20	0
6	1	0	49f3a000	20	20	0

In child with pid 6, ppid 1

PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	10	1
3	1	0	49f6a000	12	12	1
4	1	0	49f5a000	16	16	1
5	1	0	49f4a000	14	14	1
6	1	1	49f3a000	18	20	0

...

PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	10	1
3	1	0	49f6a000	12	10	4
4	1	0	49f5a000	16	10	12
5	1	0	49f4a000	14	10	8
6	1	1	49f3a000	18	11	14

PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0

Dec 09, 15 0:39			out.txt			Page 2/3
2	1	0	49f7a000	10	10	1
3	1	0	49f6a000	12	10	4
4	1	0	49f5a000	16	10	12
5	1	0	49f4a000	14	10	8
6	1	1	49f3a000	18	11	15
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	1	49f7a000	10	10	1
3	1	0	49f6a000	12	10	4
4	1	0	49f5a000	16	10	12
5	1	0	49f4a000	14	10	8
6	1	0	49f3a000	18	10	16
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	9	2
3	1	1	49f6a000	12	10	4
4	1	0	49f5a000	16	10	12
5	1	0	49f4a000	14	10	8
6	1	0	49f3a000	18	10	16
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	9	2
3	1	1	49f6a000	12	10	5
4	1	0	49f5a000	16	10	12
5	1	0	49f4a000	14	10	8
6	1	0	49f3a000	18	10	16
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	9	2
3	1	0	49f6a000	12	9	6
4	1	1	49f5a000	16	10	12
5	1	0	49f4a000	14	10	8
6	1	0	49f3a000	18	10	16
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	9	2
3	1	0	49f6a000	12	9	6
4	1	1	49f5a000	16	10	13
5	1	0	49f4a000	14	10	8
6	1	0	49f3a000	18	10	16
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	9	2
3	1	0	49f6a000	12	9	6
4	1	0	49f5a000	16	9	14
5	1	1	49f4a000	14	10	8
6	1	0	49f3a000	18	10	16
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	9	2
3	1	0	49f6a000	12	9	6
4	1	0	49f5a000	16	9	14
5	1	1	49f4a000	14	10	9
6	1	0	49f3a000	18	10	16
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	0	0
2	1	0	49f7a000	10	9	2
3	1	0	49f6a000	12	9	6
4	1	0	49f5a000	16	9	14

Dec 09, 15 0:39			out.txt			Page 3/3
5	1	0	49f4a000	14	9	10
6	1	1	49f3a000	18	10	16
...						
Child returned with code 3						
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	20	0
2	1	1	49f7a000	10	1	18
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	20	0
2	1	1	49f7a000	10	1	19
PID	PPID	STATE	STACK	STATIC	DYNAMIC	TIME
1	0	2	49f98000	20	20	0
2	1	1	49f7a000	10	0	20
Child returned with code 2						
Children returned in order: 6 4 5 3 2						
Ending init process...						